

# COMS30030 - Image Processing and Computer Vision

Lecture 05

## Segmentation - The Basics

Majid Mirmehdi  
[majid@cs.bris.ac.uk](mailto:majid@cs.bris.ac.uk)



# Examples of Image Segmentation

- **Image Segmentation ...**

... is the process of spatial subsectioning of a (digital) image into multiple partitions of pixels (i.e. segments or regions) according to given criteria.




Example: segmentation of an image into locally coherent regions

# Motivation: Why Segment Images?


- **Image Simplification**

- an image may contain millions of pixels but only a few regions



- **Higher-level Object Description**

- regions tend to belong to the same class of object
- regions may provide object properties (e.g. shape, colour, ...)






- **Input for Content Classifiers**

- region descriptions can be input data for higher level classifiers, e.g. Bayesian Classifiers or Neural Networks.



# Why Segment Images?



Examples from <https://medium.com/cogitotech> and Alberto Pretto





# Grouping Pixels

## Goals:

- Gather pixels/features that belong together
- Obtain an intermediate representation that compactly describes key image (video) parts

## Top-down vs. bottom-up segmentation

- Top-down: pixels belong together because they are from the same object
- Bottom-up: pixels belong together because they look similar



Hard to measure success: what is interesting depends on the application.

# Example of Over-Segmentation

Original image



Over-segmentation




***Over-segmentation:*** pixels belonging to the same region [object] are classified as belonging to different regions [objects]

# Example of Under-Segmentation

Original image



Under-segmentation




*Under-segmentation:* pixels belonging to different regions [objects] are classified as belonging to the same region [object]

# So many segmentation methods...


## Thresholding Methods

- pixels are categorized based on intensity
- only useful when sufficient contrast exists



## Edge-based Methods

- region boundaries are constructed from edgemaps



## Region-based Methods

- region growing from seed pixels
- region splitting and merging for efficient spatial encoding





# So many segmentation methods...


## Clustering and Statistical Methods

- global, often histogram based image partitioning, e.g. *K-means*, Gaussian Mixture Model



## Topographic Methods (out of scope in this unit)


- stepwise simplifications that take spatially wider (topographical) image configurations into account e.g. watershed transform, variational based methods




...and many more...

# Image Histogram


Brighter




Histogram




Darker



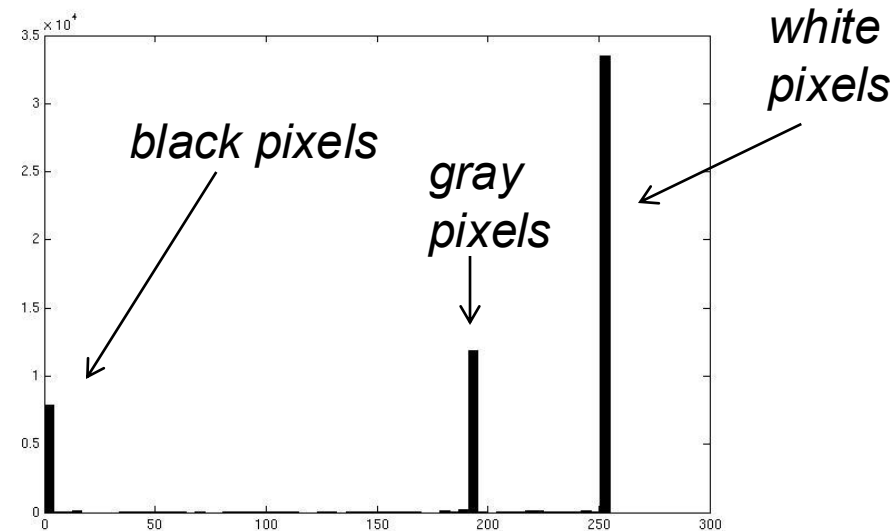
Histogram



# Image segmentation: toy example




input image




- The intensities define the three groups.
- We could label every pixel in the image according to which of these primary intensities it is.
  - i.e., *segment* the image based on the intensity feature.

# Simple thresholding not enough!

What if the image isn't quite so simple?



**input image**







# Image Segmentation

## Perfect segmentation is difficult to achieve:


- a pixel may straddle the “real” boundary of objects such that it partially belongs to two or more objects
- effects of noise, non-uniform illumination, occlusions etc. give rise to the problem of *over-segmentation* and *under-segmentation*



Images from [craftofcoding.wordpress.com](http://craftofcoding.wordpress.com)

# Thresholding Example

- If the image contains a dark object on a light background
  - choose a threshold value,  $T$
  - for each pixel
    - if the brightness at that pixel is less than  $T$ , it is a pixel of interest
    - otherwise it is part of the background
- The value of the threshold is very important
  - if too high  $\rightarrow$  background pixels classified as foreground
  - If too low  $\rightarrow$  foreground pixels classified as background




$T = 128$



$T = 96$



$T = 64$




# Using Histograms to Stipulate Regions

Maybe apply multiple thresholds?


The seal image shows three regions

- one below  $T_1 = 80$
- one above  $T_2 = 142$
- one between the two thresholds



# Iterative Threshold Selection Algorithm

1. Select an initial estimate for the threshold  $T$
2. Segment the image using  $T$ .  
This will produce two groups of pixels:  $G_1$  consisting of all pixels with grey levels  $>T$  and  $G_2$  consisting of pixels with grey values  $\leq T$ .
3. Compute the average grey level values  $m_1$  and  $m_2$  for the pixels in regions  $G_1$  and  $G_2$ .
4. Compute a new threshold value:  $T = (m_1 + m_2)/2$
5. Repeat steps (2.) through (4.) until convergence



- initial estimate
- - - average values (round 1)
- - - average values (round 2)
- threshold after round 1
- threshold after round 2

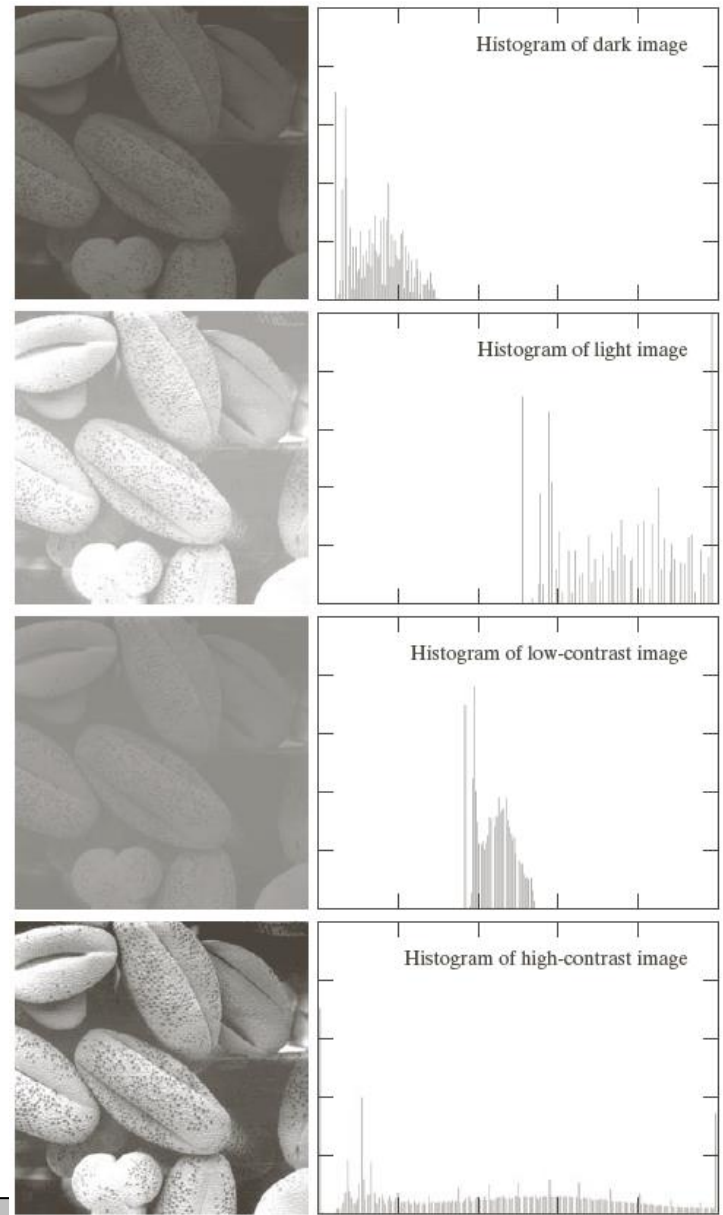




# Thresholding for Segmentation

Not always a good solution!

Four problematic image types: dark, light, low contrast, and high contrast.




# Split & Merge Segmentation – Divide & Conquer

## Homogeneity function $H$

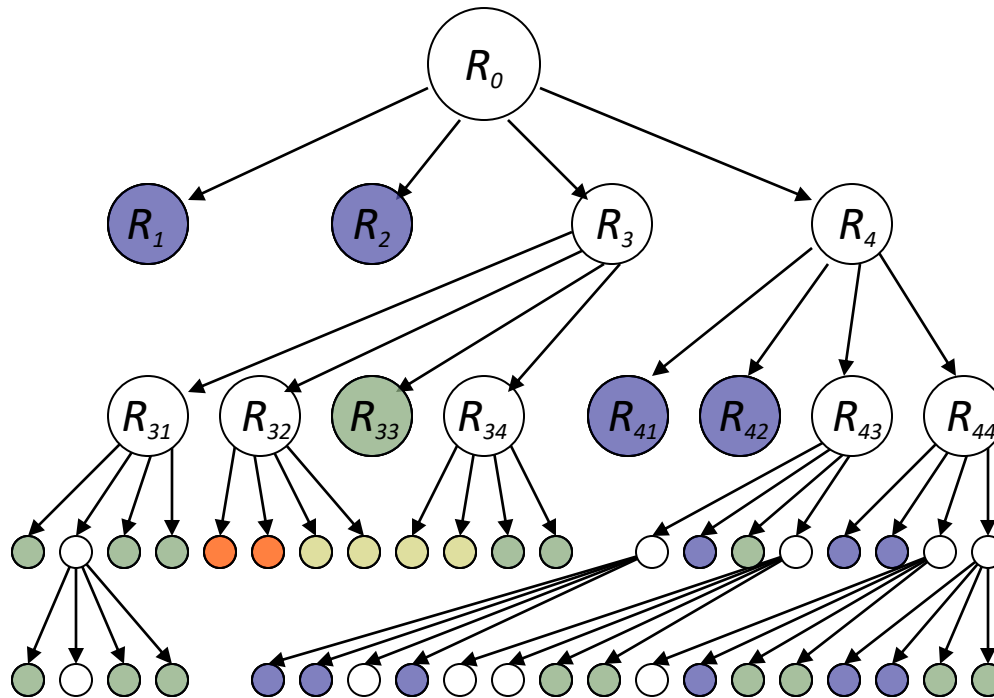
$$H(\text{Region } A) = 1 \quad (\text{homogeneous})$$

$$H(\text{Region } B) = 0 \quad (\text{inhomogeneous})$$





# Split & Merge Segmentation – Divide & Conquer

1. Start with  $R_0$  that represents the entire image
2. If  $H(R_i) = 0$  (inhomogeneous) then  
{split area into 4 blocks (quadtree splitting) and process each area with step (2.)}
3. Merge all subregions that pairwise satisfy  $H(R_i \cup R_j) = 1$  (homogenous)



Image



Segmentation Result

# Split & Merge – Summary

## Conceptual Summary:


- Iteratively decompose an image into regions of a maximally sized selected shape (e.g. rectangle) that do not satisfy a homogeneity condition. (split step)
- Then merge regions that together satisfy a homogeneity condition. (merge step)

## Some Comments:

- Using quadtrees, the results of split and merge tend to be *blocky*.
- Can have an adaptive homogeneity condition that, for instance, changes depending on the region size.

## Example H

- $H(R_i)=1$  if at least 80% of the pixels in  $R_i$  have the property  $|z_j - m_i| < 2\sigma_i$  where  $z_j$  is the grey level of the  $j^{th}$  pixel in  $R_i$ ,  $m_i$  is the mean grey level of the region and  $\sigma_i$  is the standard deviation of the grey levels in  $R_i$
- If  $H(R_i)=1$  then set all the pixels in  $R_i$  to value  $m_i$



Original

Result




# Split & Merge – Bristol Video Scene Segmentation

Original Video




Segmentation Result




- Images are segmented using a Split-And-Merge technique. (Note the blocky nature of the regions!)
- Regions are then labelled by a Neural Network to associate the segments with semantics (colouration).
- This project dates back to around 27 years ago!

# Image segmentation: toy example

What if the image isn't quite so simple?




input image




intensity

One answer is: use **clustering**...


# Clustering dilemma



If we knew the **cluster centers**, we could allocate points to groups by assigning each to its closest center.




If we knew the **group memberships**, we could get the centers by computing the mean per group.



**A “chicken and egg” problem!**

# Image segmentation: toy example



- Goal: choose three “centres” as the **representative intensities**, and label every pixel according to which of these centres it is nearest to.
- Best cluster centres are those that minimize SSD between all points and their nearest cluster centre  $\mu_j$

$$\Theta(\text{clusters}, \text{data}) = \sum_{j \in \text{clusters}} \left[ \sum_{i \in j^{\text{th}} \text{cluster}} \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2 \right]$$





# Clustering for image segmentation




map to 3D

RGB space



map back to

pixel space




# K-means clustering – theoretical view

- It minimises the following objective function:


$$\Theta(\text{clusters}, \text{data}) = \sum_{j \in \text{clusters}} \left[ \sum_{i \in j^{\text{th}} \text{ cluster}} \left\| \mathbf{x}_i - \boldsymbol{\mu}_j \right\|^2 \right]$$

- An iterative clustering algorithm
  - Pick  $K$  random points as cluster centres (means)
  - Iterate:
    - Assign data instances to closest mean
    - Assign each mean to the average of its assigned points
    - Stop when no point's assignment changes




# K-means clustering

1. Ask user how many clusters they'd like (e.g.,  $K=5$ )




# K-means clustering

1. Ask user how many clusters they'd like (e.g.,  $K=5$ )
2. Randomly guess  $K$  cluster centre locations ( $\mu_1 \dots \mu_K$ )
3. Each datapoint finds out which centre it's closest to (thus each centre "owns" a set of datapoints)




# K-means clustering

1. Ask user how many clusters they'd like (e.g.,  $K=5$ )
2. Randomly guess  $K$  cluster centre locations ( $\mu_1 \dots \mu_K$ )
3. Each datapoint finds out which centre it's closest to
4. Each centre finds the centroid of the points it owns...
5. ...and jumps there




# K-means clustering

1. Ask user how many clusters they'd like (e.g.,  $K=5$ )
2. Randomly guess  $K$  cluster centre locations ( $\mu_1 \dots \mu_K$ )
3. Each datapoint finds out which centre it's closest to
4. Each centre finds the centroid of the points it owns...
5. ...and jumps there
6. Repeat from 3 until terminated!





# K-means gone wrong!



# Reflection on the $K$ -means Algorithm

- **What does it do?**

- $K$ -means attempts to find a configuration  $\mu_1 \dots \mu_K$  that minimises within-cluster scatter: total squared distance between point  $x_i$  and centroid  $\mu_j$  in  $j^{\text{th}}$  cluster:

$$\sum_i \|\mathbf{x}_i - \boldsymbol{\mu}_j\|^2$$

- This is equivalent to maximising the between-cluster scatter (total squared distance between each cluster centroid and the global centroid of all points)

- **Does it work?**

1. The algorithm terminates.
2. It finds a local optimum from which no further improvement is possible by making local changes.
3. It does not necessarily find a global optimum.


# K-means Pros and Cons

## Pros


- Simple, fast to compute
- Converges to local minimum of within-cluster squared error

## Cons/issues


- Setting  $K$ ?
- Sensitive to initial centres
- Sensitive to outliers
- Detects spherical clusters




(A): Undesirable clusters



(B): Ideal clusters



(A): Two natural clusters



(B):  $k$ -means clusters

# Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **intensity** similarity





Feature space: intensity value (1D)

# Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **intensity**  
**AND position** similarity




Both regions are black, but if we also include **position (x,y)**, then we could group the two into distinct segments; so encode both similarity & proximity.

# Segmentation as clustering

Depending on what we choose as the *feature space*, we can group pixels in different ways.

Grouping pixels based on **colour** similarity



$$\begin{pmatrix} R=255 \\ G=200 \\ B=250 \end{pmatrix}$$

$$\begin{pmatrix} R=245 \\ G=220 \\ B=248 \end{pmatrix}$$

$$\begin{pmatrix} R=15 \\ G=189 \\ B=2 \end{pmatrix}$$

$$\begin{pmatrix} R=3 \\ G=12 \\ B=2 \end{pmatrix}$$

Slide inspired from Kristen Grauman

Feature space: intensity values (3D)



# K-means Colour Segmentation

Original image



$K = 2$



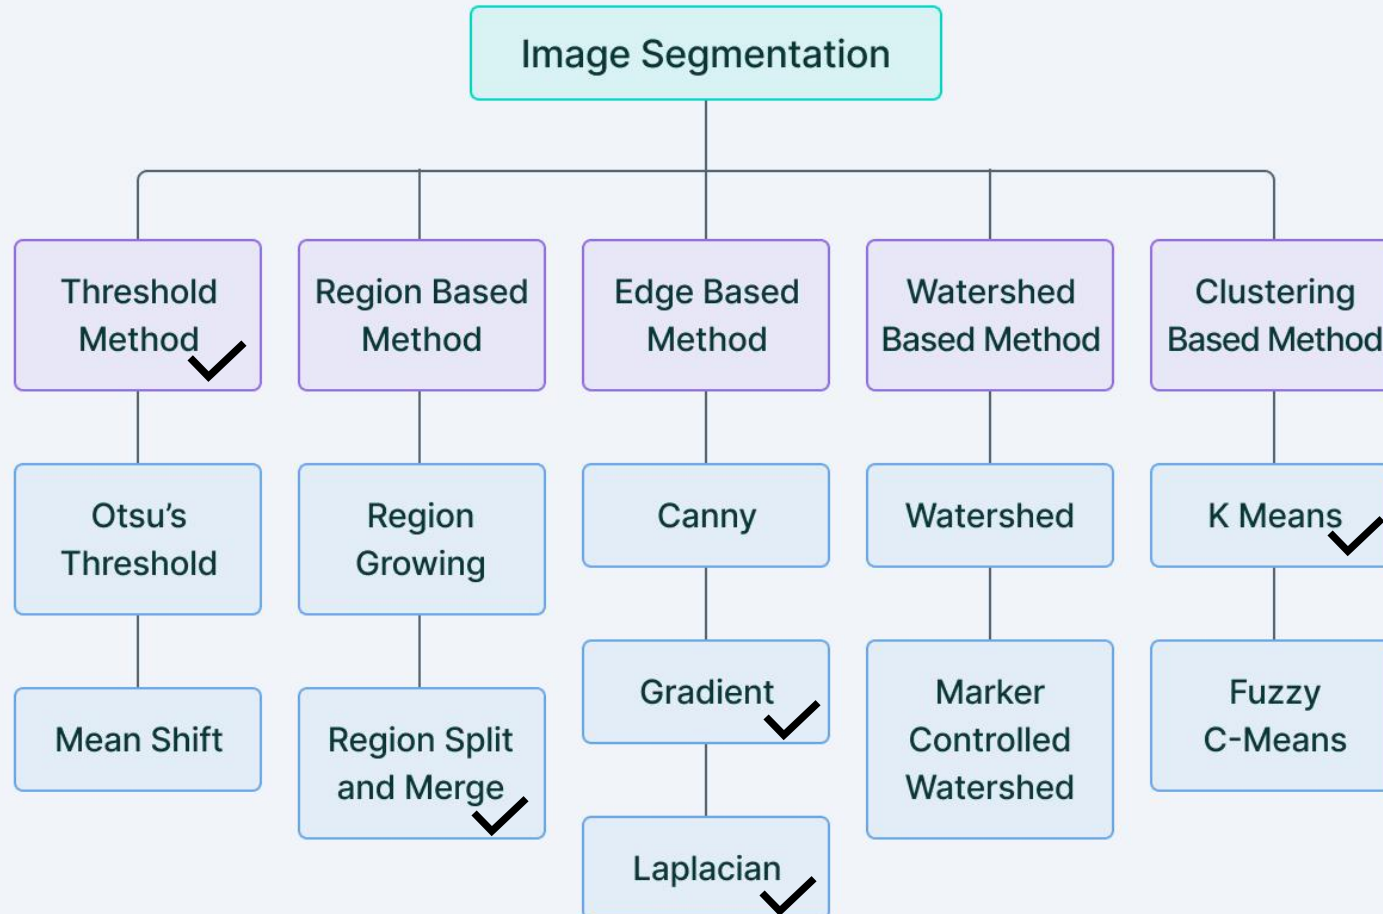
$K = 3$



$K = 10$



# Summary: Image Segmentation



V7 Labs

# Latest Segmentation Research


## Semantic Segmentation vs. Instance Segmentation vs. Panoptic Segmentation



(a) Image



(b) Semantic Segmentation



(c) Instance Segmentation



(d) Panoptic Segmentation

V7 Labs

## Object Detection